

OpenMP Prefix Sum and MPI Finding an element in a large array

Harsh Kumar, 24576

21 November 2024

1. Q1

Ans: We have solved the prefix problem for integer arrays of sizes 10,000, 20,000, and 30,000 elements using different thread counts (2, 4, 8, 16, 32, and 64). To ensure accuracy and consistency, we ran the code five times for each configuration and recorded the average execution time. We also calculated speedups relative to sequential execution and visualized the results through graphs.

Methodology for Parallel Program

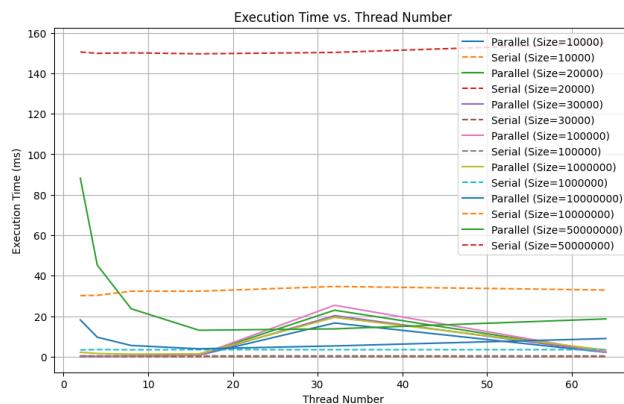
- **Array Division:** The input array A[] is divided into blocks. Each block is assigned to a different OpenMP thread based on its thread ID.
- **Local Computation:** Each thread computes the prefix sum of its assigned block. Specifically, the B[] array for each thread is calculated, with the result of each block's last element stored as a partial sum.
- **Global Adjustment:** After all threads compute their local prefix sums, an additional step is performed by a single thread to adjust the prefix sums for threads based on the partial sums. This step ensures that the results from different threads are correctly combined.
- **Final Calculation:** The prefix sum from the partial sums array is then propagated across the entire array B[] using multiple threads.

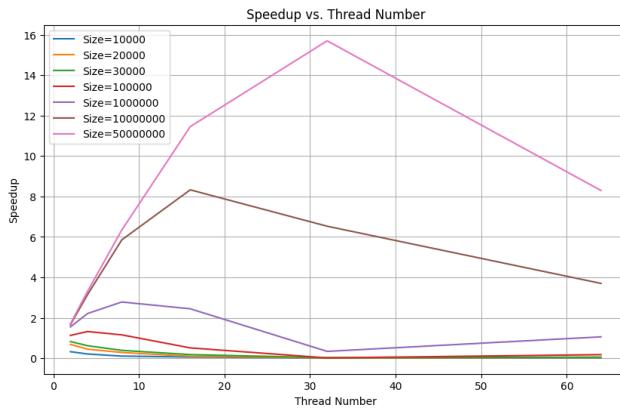
Experiments :

The program was executed for various combinations of:

- **Array sizes:** 10000, 20000, 30000, 100000, 1000000, 10000000, 50000000. (Tested on extra sizes for showcasing speed up, as the size given in question were not providing good speedup)
- **Thread counts:** 2, 4, 8, 16, 32, 64.
- **Repetitions:** Each configuration was run 5 times, and the average execution time was recorded.

Results





Data

Size	Thread Num	Parallel Time (ms)	Serial Time (ms)	Speedup
10000	2	0.127304	0.040490	0.325840
10000	4	0.201487	0.042146	0.210099
10000	8	0.405309	0.041840	0.106444
10000	16	0.757608	0.045308	0.059818
10000	32	16.631042	0.042970	0.016695
10000	64	2.130471	0.044910	0.021094
20000	2	0.119120	0.079314	0.684536
20000	4	0.190678	0.085230	0.451163
20000	8	0.283442	0.078167	0.282820
20000	16	0.863872	0.082994	0.099077
20000	32	22.950946	0.083362	0.018574
20000	64	2.269389	0.083534	0.037170
30000	2	0.144250	0.117286	0.825022
30000	4	0.195026	0.119916	0.619767
30000	8	0.329071	0.123604	0.389613
30000	16	0.701772	0.123556	0.180435
30000	32	20.271132	0.124114	0.025076
30000	64	2.275345	0.125834	0.055445

Table 1: Performance Metrics for Different Problem Sizes and Thread Numbers (given in question)

Size	Thread Num	Parallel Time (ms)	Serial Time (ms)	Speedup
50000000	8	23.674024	150.089310	6.340838
50000000	16	13.061660	149.592113	11.453240
50000000	32	13.745516	150.292523	15.697788
50000000	64	18.663620	154.784480	8.296378
10000000	4	9.614381	30.302174	3.151940
10000000	8	5.531018	32.336211	5.855866
10000000	16	3.888767	32.343829	8.328022
10000000	32	5.315737	34.679534	6.528956

Table 2: Best Speedups for Different Problem Sizes and Thread Numbers

For the complete table and data please refer to the third section.

Observations:

- Small Problem Sizes: Parallelism is inefficient for small arrays, as threading overhead outweighs computation benefits, leading to poor or negative speedups.

- Large Problem Sizes: Speedup improves significantly for larger arrays with higher thread counts, but again diminishes due to overhead and contention.
- Best Performance: Maximum speedup (15.7x) is achieved for the largest problem size (50,000,000 elements)
- Excessive thread counts (32–64) generally degrade performance for our problem due to synchronization overhead.

2. Q2

Ans: We have implemented an MPI-based parallel search algorithm that can find a target element in a distributed array across multiple processes. The steps are as follows:

1. Initialization:

- We initialize MPI and obtain the rank and size of the processes.
- Root process (`rank 0`) generates a random integer array of size 1,000,000 with values between 1 and 5,000,000.

2. Array Distribution:

- The array is distributed across processes using `MPI_Scatter`.

3. Broadcast Target:

- A target integer is broadcast to all processes using `MPI_Bcast`.

4. Local Search:

- Each process searches its local array for the target element.
- If found, the global index is calculated and other processes are notified via `MPI_Send`.

5. Termination Notification:

- Processes use `MPI_Irecv` to listen for termination signals and stop searching upon receipt.
- Processes does not locate the target immediately, it checks for messages from other processes using the `MPI_Test` function after each 1000 elements (to reduce communication overhead).
 - If received, Update the global find and exit the loop

6. Result Aggregation:

- The global index of the target is collected using `MPI_Allreduce` and displayed by the root process.

7. Performance Measurement:

- Execution times are recorded for configurations of 1, 8, 16, 32, and 64 processes.
- 20 random targets are tested for each configuration to calculate average execution time.

8. Cleanup:

- Memory is cleaned up and MPI is finalized.

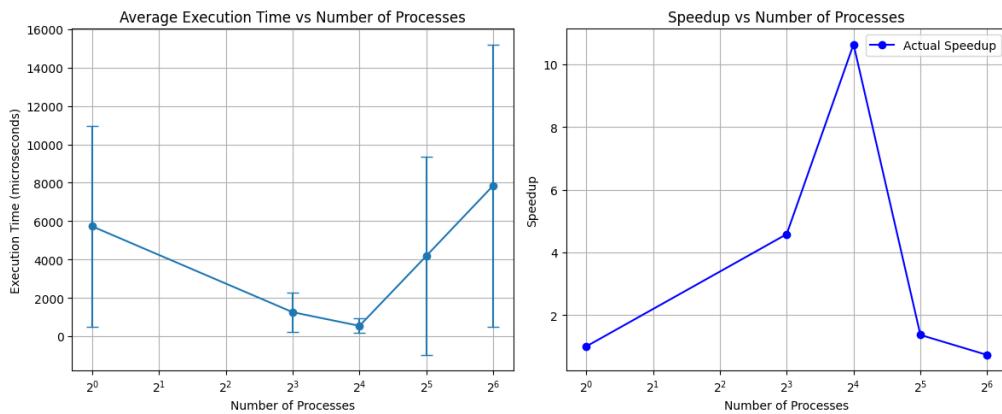
Results

Execution times and speedup for various numbers of processes are summarized below:

Processes	Avg Time (μ s)	Std Dev	Speedup
1	5726.90	5251.40	1.00
8	1250.00	1015.23	4.58
16	538.55	382.43	10.63
32	4184.90	5165.66	1.37
64	7853.45	7360.73	0.73

Table 3: Performance Summary

Speedup graphs



Observations

- The program achieves significant speedup up to 16 processes due to effective load distribution.
- Beyond 16 processes, communication overhead reduces the benefits of parallelism.

Conclusion

The MPI-based search program demonstrates scalability for up to 16 processes, showcasing efficient parallelism within the limits of communication overhead.

3. Complete Data for Question 1

Ans: PTO

	size	thread_num	parallel_time	serial_time	speedup
0	10000	2	0.127304	0.040490	0.325840
1	10000	4	0.201487	0.042146	0.210099
2	10000	8	0.405309	0.041840	0.106444
3	10000	16	0.757608	0.045308	0.059818
4	10000	32	16.631042	0.042970	0.016695
5	10000	64	2.130471	0.044910	0.021094
6	20000	2	0.119120	0.079314	0.684536
7	20000	4	0.190678	0.085230	0.451163
8	20000	8	0.283442	0.078167	0.282820
9	20000	16	0.863872	0.082994	0.099077
10	20000	32	22.950946	0.083362	0.018574
11	20000	64	2.269389	0.083534	0.037170
12	30000	2	0.144250	0.117286	0.825022
13	30000	4	0.195026	0.119916	0.619767
14	30000	8	0.329071	0.123604	0.389613
15	30000	16	0.701772	0.123556	0.180435

16	30000	32	20.271132	0.124114	0.025076
17	30000	64	2.275345	0.125834	0.055445
18	100000	2	0.375477	0.405741	1.130139
19	100000	4	0.324188	0.419853	1.321134
20	100000	8	0.366781	0.423689	1.156880
21	100000	16	0.812724	0.410131	0.512629
22	100000	32	25.389135	0.412759	0.016455
23	100000	64	2.228281	0.398051	0.178756
24	1000000	2	2.137081	3.310159	1.550438
25	1000000	4	1.600611	3.509708	2.212294
26	1000000	8	1.242307	3.408323	2.782206
27	1000000	16	1.427149	3.475516	2.445694
28	1000000	32	19.400570	3.441069	0.339733
29	1000000	64	3.303789	3.469768	1.058547
30	10000000	2	18.165632	30.152856	1.660598
31	10000000	4	9.614381	30.302174	3.151940
32	10000000	8	5.531018	32.336211	5.855866
33	10000000	16	3.888767	32.343829	8.328022
34	10000000	32	5.315737	34.679534	6.528956

35	10000000	64	8.950479	32.937430	3.701584
36	50000000	2	88.155199	150.507101	1.707388
37	50000000	4	45.229633	149.824115	3.312820
38	50000000	8	23.674024	150.089310	6.340838
39	50000000	16	13.061660	149.592113	11.453240
40	50000000	32	13.745516	150.292523	15.697788
41	50000000	64	18.663620	154.784480	8.296378